

# BHAVUK AGRAWAL

bhavukagrawal1409@gmail.com | +91-7060802357 | github.com/bhavuk1409 | linkedin.com/in/bhavukagrawal1409

## EDUCATION

**Bennett University**, Greater Noida, Uttar Pradesh, India  
Bachelor of Technology in Computer Science | **CGPA: 9.0/10**

**Aug 2023 – 2027**

## EXPERIENCE

**Machine Learning Intern** | *National University of Singapore (NUS) Singapore (On-Site)*

**June 2024 – July 2024**

- Developed TF-IDF-based medical report analysis tool to accelerate text processing workflows for the research team
- Automated deployment infrastructure using AWS CloudFormation, reducing manual setup effort
- Deployed ML models on AWS SageMaker and optimized pipeline execution using AWS services
- Collaborated with researchers to transform academic ML/NLP concepts into functional prototypes and demos

## PROJECTS

**R-Search: Reinforcement Learning for Multi-Source Search** | *Python, PyTorch, LoRA, HuggingFace, WandB*

github.com/bhavuk1409/llm-reasoning-search

- Developed novel GRPO-based RL fine-tuning pipeline for LLM reasoning with structured output generation, integrating Exa multi-source search API and LoRA parameter-efficient training with HuggingFace and WandB support
- Designed composite reward function combining semantic coherence, DAG validation, and format compliance with curriculum learning, optimizing training stability and convergence
- Achieved 25-30% accuracy improvement on HotpotQA-hard benchmark, 40% better format compliance, and 35% VRAM reduction through LoRA optimization, enabling training on consumer GPUs

**LLM-Enhanced RL Trading System** | *Python, PPO, LangChain, OpenRouter, Streamlit, Exa API*

github.com/bhavuk1409/llm-rl-trading

- Engineered hybrid PPO + multi-agent LLM trading framework integrating real-time sentiment via Exa API, delivering 18-25% higher returns vs baseline PPO and 40% lower action volatility
- Implemented coordinated technical, fundamental, sentiment, and risk agents using LangChain and OpenRouter with interactive Streamlit dashboard for strategy analysis

**AI-Powered Real-Time Video Conferencing Platform** | *WebRTC, AWS, Whisper, BART, LLaMA, LangChain, RAG*

github.com/bhavuk1409/intellimeet\_updated

- Architected scalable platform with real-time multilingual subtitles (Whisper), auto-summarization (BART), and RAG chatbot (LLaMA + LangChain) on AWS-compatible infrastructure with WebRTC for low-latency communication

## TECHNICAL SKILLS

**Languages:** Python, C++, Java, JavaScript, SQL

**ML/AI Frameworks:** PyTorch, TensorFlow, Transformers, LLM Fine-tuning (LoRA/QLoRA/PEFT), Reinforcement Learning (PPO, GRPO)

**AI/LLM Tools:** LangChain, LangGraph, RAG, Hugging Face, OpenAI API, Vector Databases, Whisper, BART

**Cloud & DevOps:** AWS (SageMaker, CloudFormation, EC2, S3), Docker, Git

**Web Development:** FastAPI, React, MongoDB, WebRTC, Streamlit

**Core CS:** Data Structures & Algorithms, System Design, Database Management, Problem Solving

## ACHIEVEMENTS & CERTIFICATIONS

- Selected for competitive on-site ML/AI research internship at National University of Singapore (acceptance rate <5%)
- Solved 400+ Data Structures & Algorithms problems across coding platforms | **LeetCode Rating: 1644** (Top 25%)
- Active contributor to open-source AI projects with focus on LLM reasoning and reinforcement learning applications
- **NVIDIA:** Getting Started with Accelerated Computing in Modern CUDA C++
- **DeepLearning.AI:** Improving Deep Neural Networks – Hyperparameter Tuning, Regularization and Optimization
- **UC San Diego:** Data Structures Specialization